

# Minimax Implementation

## Recursive Approach

def max-value(state):


  initialize  $v = -\infty$

  for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

  return  $v$

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



def min-value(state):

  initialize  $v = +\infty$

  for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

  return  $v$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

# Minimax Implementation (Dispatch)

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

def max-value(state):

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

def min-value(state):

initialize  $v = +\infty$


for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$




return v

# Minimax Efficiency

- How efficient is minimax?
  - Complete: if tree is finite
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$ (why?)
  - Space:  $O(bm)$ (why?)
- Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?



# Minimax Properties




Optimal against a perfect player. Otherwise?

[Demo: min vs exp (L6D2, L6D3)]

# Video of Demo Min vs. Exp (Min)

---


Ghosts are intelligence



# Video of Demo Min vs. Exp (Exp)

---

Ghost are (some time) in Random Action




# Resource Limits

---



# Challenge: Resource Limits

- Problem: In realistic games, cannot search to leaves!  
Usually search is Time limited!!
- Solution: Depth-limited search
  - Instead, search only to a limited depth in the tree (is it Ok?)
  - Replace terminal utilities with an **evaluation function** for non-terminal positions
- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\alpha$ - $\beta$  reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference()
- Use iterative deepening for an anytime algorithm(if we have time search in new depth)






# Video of Demo Limited Depth (2) **Low value**

---


What happen if we  
search just 2 level



# Video of Demo Limited Depth (10) Good value


---

What happen if we  
search all state




# Evaluation Functions

---



# Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search




- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.


# Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the **tradeoff** between **complexity of features** and **complexity of computation**



[Demo: depth limited (L6D4, L6D5)]


# Evaluation for Pacman




[Demo: thrashing  $d=2$ , thrashing  $d=2$  (fixed evaluation function), smart ghosts coordinate (L6D6,7,8,10)]

# Video of Demo Thrashing (d=2)

---




# Why Pacman Starves



- assume
  - evaluation function  $\rightarrow 10 * \text{number of eaten dots}$
  - Search in depth 2
  - What is Value of Root
  - If select right
    - Start New Search
    - what is value of root? (What happen If select left)
- A danger of replanning agents!
  - He knows his score will go up by eating the dot now (west, east)
  - He knows his score will go up just as much by eating the dot later (east, west)
  - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
  - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!



# Why Pacman Starves



- **assume**


- evaluation function  $\rightarrow 10 * \text{number of eaten dots}$
- Search in depth 2
- What is Value of Root
- If select right
  - Start New Search
  - what is value of root? (What happen If select left)

- **Solution**

- Distance to dot


# Video of Demo Thrashing -- Fixed ( $d=2$ )

---




# Video of Demo Smart Ghosts (Coordination)

---




# Video of Demo Smart Ghosts (Coordination) – Zoomed In

---




# Game Tree Pruning

---



# Minimax Example

1. Min Search
  1. Select 3
    1. First 3
    2. Second 12
    3. Third 8
  2. Select 2
    1. First 2
    2. Second 4
    3. Third 6
  3. Select 2
  - ...
2. Max Search
  1. Select 3



Is it efficient (computational Cost) ? Why?

# Alpha-Beta Example

## 1. Min Search

### 1. Select 3

1. First 3
2. Second 12
3. Third 8

### 2. Select 2

1. First 2
2. Second 4 **prune**
3. Third 6


### 3. Select 2

...

## 2. Max Search

### 1. Select 3


$\alpha$  = best option so far from any MAX node on this path



- ***The order of generation matters***: more pruning is possible if good moves come first

# Alpha-Beta Pruning

- General case (pruning children of MIN node)
  - We're computing the MIN-VALUE at some node  $n$
  - We're looping over  $n$ 's children
  - $n$ 's estimate of the childrens' min is dropping
  - Who cares about  $n$ 's value? MAX
  - Let  $\alpha$  be the best value that MAX can get so far at any choice point along the current path from the root
  - If  $n$  becomes worse than  $\alpha$ , MAX will avoid it, so we can prune  $n$ 's other children (it's already bad enough that it won't be played)
- Pruning children of MAX node is symmetric
  - Let  $\beta$  be the best value that MIN can get so far at any choice point along the current path from the root





# Alpha-Beta Pruning

---

- There are two rules for terminating search:
  - Search can be stopped **below any MIN** node having a **beta value less** than or equal to the **alpha** value of any of its MAX ancestors.
  - Search can be stopped **below any MAX** node having an **alpha value greater** than or equal to the **beta** value of any of its MIN ancestors.

# Alpha-Beta Implementation


$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```


# Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong (ambiguity in action selection!)
  - Important: children of the root may have the wrong value
  - So the most naïve version won't let you do action selection
  - How solve?
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth!
  - Full search of, e.g. chess, is still hopeless...
- This is a simple example of **meta reasoning** (computing about what to compute)




# Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
  - Important: children of the root may have the wrong value
  - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
  - Time complexity drops to  $O(b^{m/2})$
  - Doubles solvable depth!
  - Full search of, e.g. chess, is still hopeless...
- This is a simple example of **metareasoning** (computing about what to compute)




# Alpha-Beta Quiz


---



# Alpha-Beta Quiz 2



## Alpha-Beta Quiz 2



# Alpha-Beta algorithm

**function** ALPHA-BETA-DECISION(*state*) **returns** an action  
    **return** the *a* in ACTIONS(*state*) maximizing MIN-VALUE(RESULT(*a*, *state*))

---

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **inputs:** *state*, current state in game  
             $\alpha$ , the value of the best alternative for MAX along the path to *state*  
             $\beta$ , the value of the best alternative for MIN along the path to *state*  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow -\infty$   
    **for** *a*, *s* in SUCCESSORS(*state*) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$   
        **if**  $v \geq \beta$  **then return** *v*  
         $\alpha \leftarrow \text{MAX}(\alpha, v)$   
    **return** *v*

---

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    same as MAX-VALUE but with roles of  $\alpha$ ,  $\beta$  reversed




# Alpha-Beta algorithm




```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

# Games with uncertain outcomes




# Worst-Case vs. Average Case



Idea: Uncertain outcomes controlled by chance, not an adversary!  
Note: Change MIN node shape

# Expectimax Search

- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- **Expectimax search**: compute the average score under optimal play
  - Max nodes as in minimax search
  - Chance nodes are like min nodes but the outcome is uncertain
  - Calculate their **expected utilities**
  - I.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as **Markov Decision Processes**



# Expectimax Pseudocode

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is EXP: return exp-value(state)

```
def max-value(state):
```

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return  $v$

```
def exp-value(state):
```

initialize  $v = 0$

for each successor of state:


$p = \text{probability}(\text{successor})$

$v += p * \text{value}(\text{successor})$

return  $v$


# Expectimax Pseudocode

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```





$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

# Expectimax Example



# Expectimax Pruning?



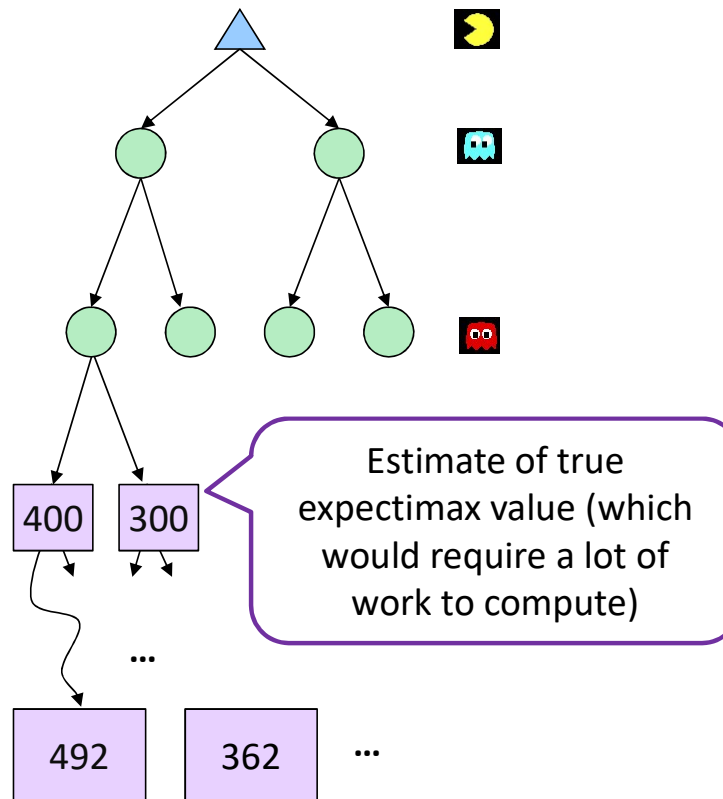
Generally Expectimax Pruning is not true.

But good, if There is upper bound for pruning node

compute upper bound for Expectation and prune based on this value




# Depth-Limited Expectimax?



# Probabilities

---



# Reminder: Probabilities

- A **random variable** represents an event whose outcome is unknown
- A **probability distribution** is an assignment of weights to outcomes
- **Example: Traffic on freeway**
  - Random variable:  $T$  = whether there's traffic
  - Outcomes:  $T$  in {none, light, heavy}
  - Distribution:  $P(T=\text{none}) = 0.25$ ,  $P(T=\text{light}) = 0.50$ ,  $P(T=\text{heavy}) = 0.25$
- **Some laws of probability (more later):**
  - Probabilities are always non-negative
  - Probabilities over all possible outcomes sum to one
- **As we get more evidence, probabilities may change:**
  - $P(T=\text{heavy}) = 0.25$ ,  $P(T=\text{heavy} \mid \text{Hour}=8\text{am}) = 0.60$
  - We'll talk about methods for reasoning and updating probabilities later



0.25



0.50






0.25

# Reminder: Expectations

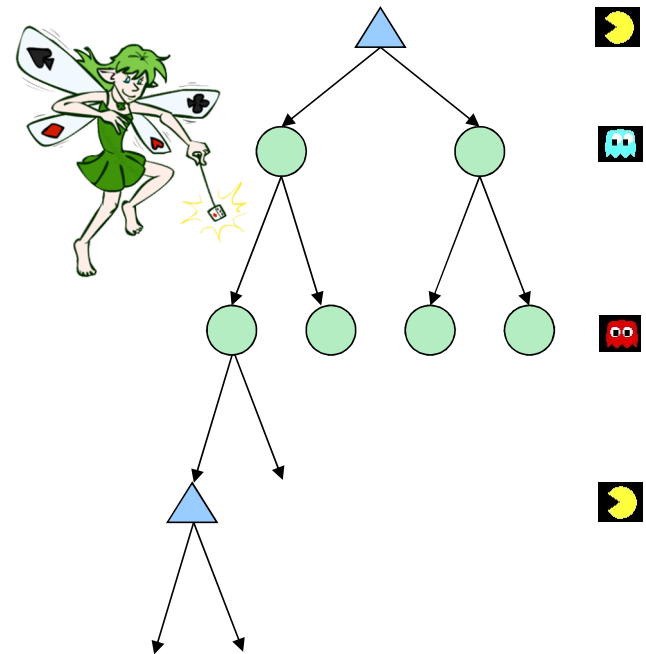
- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?

Time:	20 min		30 min		60 min		
	x		x		x		
Probability:	0.25	+	0.50	+	0.25		35 min



# What Probabilities to Use?


- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
  - Model could be a simple uniform distribution (roll a die)
  - Model could be sophisticated and require a great deal of computation
  - We have a chance node for any outcome out of our control: opponent or environment
  - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



*Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!*

# Informed Probabilities


- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



- Answer: Expectimax!
  - To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
  - This kind of thing gets very slow very quickly
  - Even worse if you have to simulate your opponent simulating you...
  - ... except for minimax, which has the nice property that it all collapses into one game tree

# Modeling Assumptions

---



# The Dangers of Optimism and Pessimism

## Dangerous Optimism

Assuming **(random) chance** when the world is adversarial




## Dangerous Pessimism

Assuming **the worst case** when it's not likely





# Assumptions vs. Reality



Pacman used depth 4 search with an eval function that avoids trouble  
Ghost used depth 2 search with an eval function that seeks Pacman

	Adversarial Ghost	Random Ghost
Minimax Pacman	Won 5/5 Avg. Score: 483	Won 5/5 Avg. Score: 493
Expectimax Pacman	Won 1/5 Avg. Score: -303	Won 5/5 Avg. Score: 503

Results from playing 5 games

Notice: Use Expectimax if you are sure the ghost is Random


70

[Demos: world assumptions (L7D3,4,5,6)]

# Video of Demo World Assumptions


## Random Ghost – Expectimax Pacman

---




# Video of Demo World Assumptions Adversarial Ghost – Minimax Pacman

---



# Video of Demo World Assumptions Adversarial Ghost – Expectimax Pacman


---



# Video of Demo World Assumptions

## Random Ghost – Minimax Pacman

---



# Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?

- Generalization of minimax:

- Terminals have utility tuples
- Node values are also utility tuples
- Each player maximizes its own component
- Can give rise to cooperation and competition dynamically...

